

Deploying Stateless Application with Deployment Objects.

Kubectl - kubectl controls the Kubernetes cluster manager

You can run an application by creating a Kubernetes Deployment object, and you can describe a Deployment in a YAML file. For example, this YAML file describes a Deployment that runs the **<your-docker-hub-username>/ril:v1** Docker image

Replace **<your-docker-hub-username>/ril:v1** with the image name, repository name and the tag from your **docker hub** account that you pushed to docker hub in **docker lab 8**.

1. Login to your AWS Instance and make a dir /home/<your-user-name>/application

```
$ cd /home/<your-username>
```

```
$ mkdir application
```

```
$ cd application/
```

```
$ vim <your-name>-deployment.yaml # paste the below text in the vim editor
```

#Note : press 'i' to start the edit mode in the vim editor.

```
apiVersion: apps/v1
```

```
kind: Deployment
```

```
metadata:
```

```
  name: <your-name>-deployment
```

```
spec:
```

```
  selector:
```

```
    matchLabels:
```

```
      app: <your-app-name>
```

```
  replicas: 2 # tells deployment to run 2 pods matching the template
```

```
  template:
```

```
    metadata:
```

```
      labels:
```

```
        app: <your-app-name>
```

```
    spec:
```

```
      containers:
```

```
        - name: <your-container-name>
```

```
          image: <docker-hub-image>/image:tag > #ex : asyed755/ril:v1 - it should be the same as on docker hub
```

```
          ports:
```

```
            - containerPort: 80
```

2. Create a Deployment based on the YAML file:

\$ kubectl apply -f <your-name>-deployment.yaml

3. Display information about the Deployment:

\$ kubectl describe deployment deployment

The output is similar to this:

```
user@computer:~/website$ kubectl describe deployment ril-deployment
Name:      <your-name>-deployment
Namespace: default
CreationTimestamp: Tue, 30 Aug 2016 18:11:37 -0700
Labels:    app=ril
Annotations: deployment.kubernetes.io/revision=1
Selector:  app=ril
Replicas:  2 desired | 2 updated | 2 total | 2 available | 0 unavailable
StrategyType: RollingUpdate
MinReadySeconds: 0
RollingUpdateStrategy: 1 max unavailable, 1 max surge
Pod Template:
  Labels:  app=ril
  Containers:
    nginx:
      Image:      asyed755/ril:v1
      Port:       80/TCP
      Environment: <none>
      Mounts:      <none>
      Volumes:     <none>
  Conditions:
    Type Status Reason
    ----
    Available True MinimumReplicasAvailable
    Progressing True NewReplicaSetAvailable
    OldReplicaSets: <none>
    NewReplicaSet:  nginx-deployment-1771418926 (2/2 replicas created)
    No events.
```

4. List the pods created by the deployment:

```
$ kubectl get pods -l app=<your-app-name>
```

The output is similar to this:

<i>NAME</i>	<i>READY</i>	<i>STATUS</i>	<i>RESTARTS</i>	<i>AGE</i>
<i>ril-deployment-1471416983-7o5ac</i>	<i>1/1</i>	<i>Running</i>	<i>0</i>	<i>16h</i>
<i>ril-deployment-1541148254-318ad</i>	<i>1/1</i>	<i>Running</i>	<i>0</i>	<i>16h</i>

5. To display information about a pod:

```
$ kubectl describe pod <pod-name>
```

6. Expose the Deployment with the below command.

```
kubectl expose deployment <your-app-name> --type=LoadBalancer --name=<your-service-name>
```

7. Login to the **Kubernetes Dashboard** and goto Services, click on the Load-Balancer endpoint associated with your Service to access the application webpage.